


# Switching gears: neural networks


---



Source: <https://thesharperdev.com/build-your-first-neural-network-part-2/>

# Why neural networks?


---





# Why neural networks?


- Best performing algorithms for complex tasks, bar none.
- Known potential of hierarchical organisation of simple units because of biological examples (though neural networks are not good models of actual neurons)
- Observation in frogs and cat visual cortex: there are specific layers of neurons, where earlier layers detect basic shapes (lines, edges) with later layers incorporating this information into more complex features about what is seen.



Source:  
[https://en.wikipedia.org/wiki/Axon#/media/File:Blauen\\_0657\\_MultipolarNeuron.png](https://en.wikipedia.org/wiki/Axon#/media/File:Blauen_0657_MultipolarNeuron.png)

# Why neural networks?

- Until now, we decided on the features to give to our algorithms: think tumour size, biopsy test scores, etc.
- With neural networks and images, the situation changes: we don't arduously describe what is in each image, but rather let the network learn to extract and combine features *so that* it can classify training examples correctly.

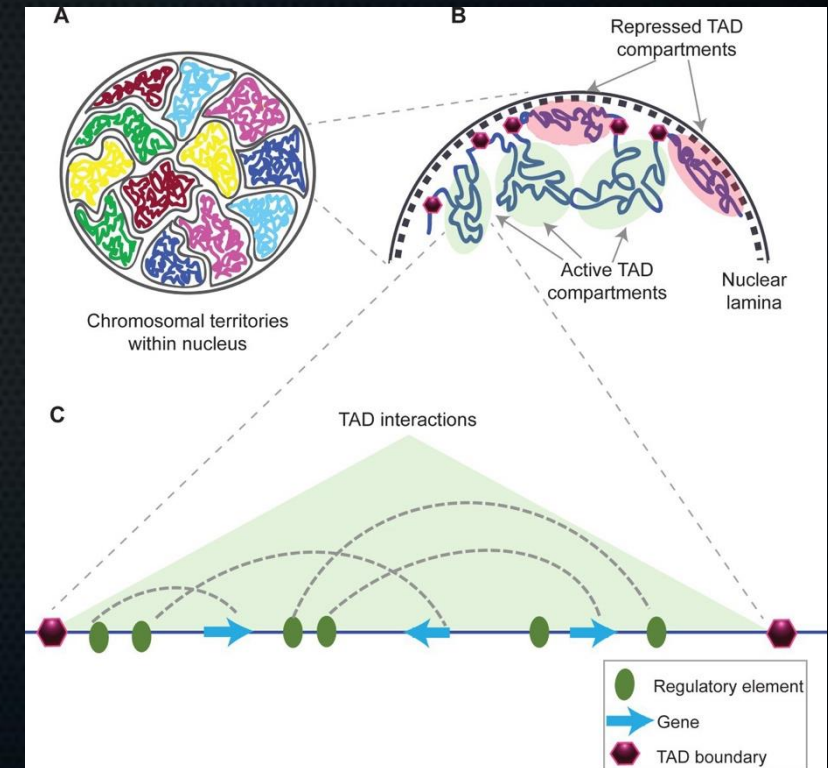





# Why neural networks?

- Caveat for biology: it can be quite difficult to translate biological problems into a framework fit for deep learning.

- Example: in images, nearby pixels probably hold similar information, i.e. are involved in the same thing. Due to (long-range) 3D-folding of DNA, linearly far DNA can be close together functionally. You need to encode your network or input to accommodate this!



Source:  
[https://en.wikipedia.org/wiki/Topologically\\_associating\\_domain#/media/File:Structural\\_organization\\_of\\_chromatin.png](https://en.wikipedia.org/wiki/Topologically_associating_domain#/media/File:Structural_organization_of_chromatin.png)

# Why neural networks?

---

- The mythical property of *universal approximation*. This says that neural networks can approximate any function with arbitrary accuracy, even with only 1 hidden layer (given enough neurons in it).

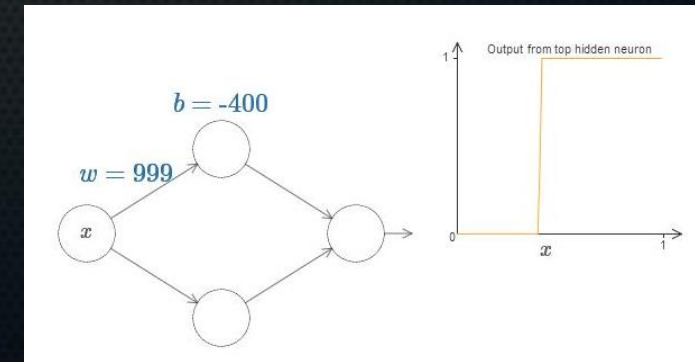

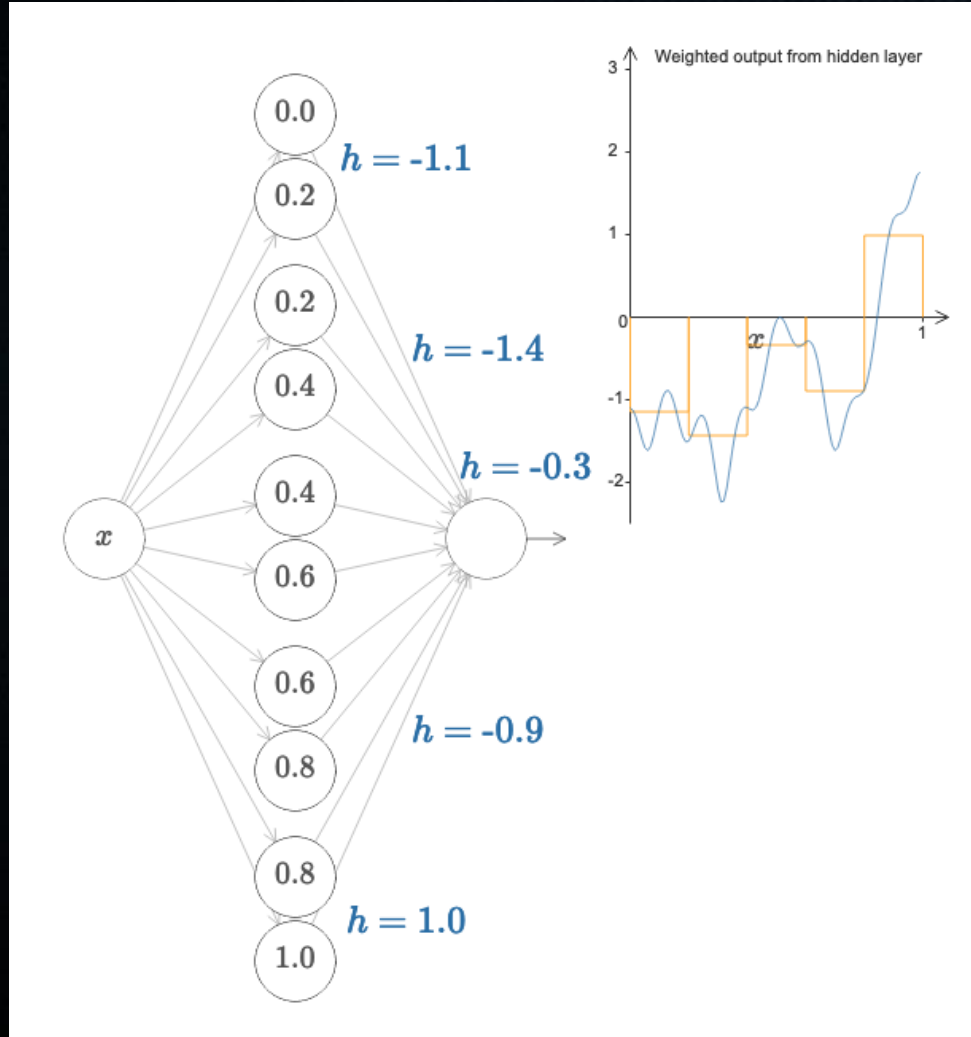


# Why neural networks?

---

- The mythical property of *universal approximation*. This says that neural networks can approximate any function with arbitrary accuracy, even with only 1 hidden layer (given enough neurons in it).
- Of course, that doesn't necessarily mean we would have the data to *train* such a neural network efficiently. Just that it is provable that for *any continuous function* a neural network can exist that approximates it as well as you like.


# Why neural networks?






# Like biology? **No**

Biological Neurons:  
Complex connectivity patterns



This image is CC0 Public Domain

Neurons in a neural network:  
Organized into regular layers for  
computational efficiency



Source : [http://cs231n.stanford.edu/slides/2019/cs231n\\_2019\\_lecture04.pdf](http://cs231n.stanford.edu/slides/2019/cs231n_2019_lecture04.pdf)

# Like biology? **No**

---

## Biological Neurons:

- Many different types
- Dendrites can perform complex non-linear computations
- Synapses are not a single weight but a complex non-linear dynamical system

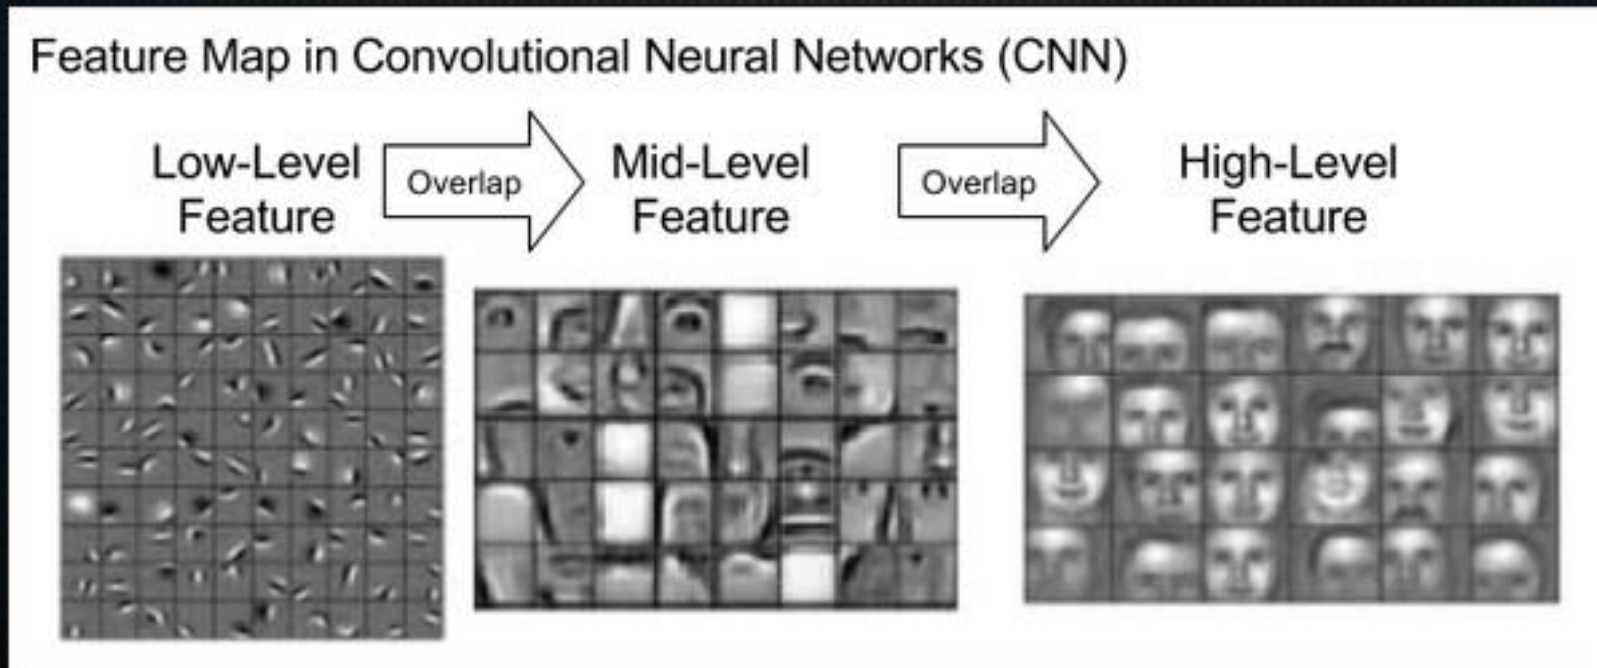
Source : [http://cs231n.stanford.edu/slides/2019/cs231n\\_2019\\_lecture04.pdf](http://cs231n.stanford.edu/slides/2019/cs231n_2019_lecture04.pdf)

- Human brains ~a cool 86 billion neurons
- A neuron can have 400.000 dendrites
- Real brains *vastly outclass* their computational analogues




# Like biology? **No**

- Still extremely useful
- Parts of how they learn *superficially resemble* how we learn



# What do neural nets have to do with logistic regression?

- A single neuron is a logistic regressor!\*




Source: <https://thedatafrog.com/en/articles/logistic-regression/>

\*Given sigmoid activation function (which is not current standard practice anymore, but was at first)




# What do neural nets have to do with logistic regression?

- A single neuron is a logistic regressor!



# What do neural nets have to do with logistic regression?


- A single neuron is a logistic regressor!





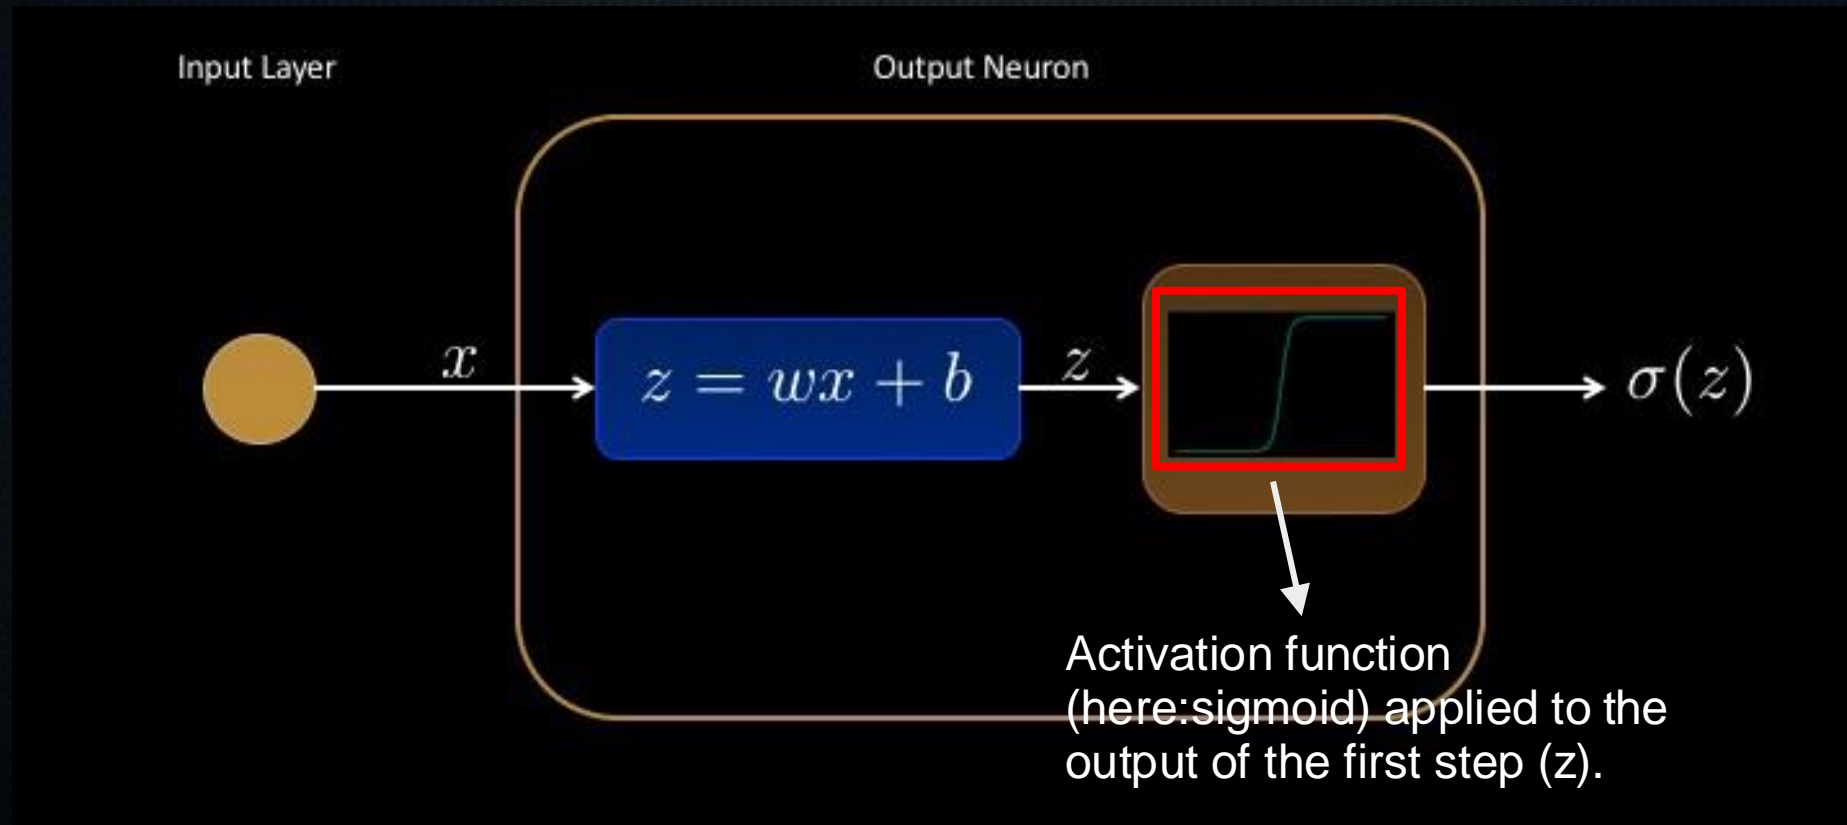
# What do neural nets have to do with logistic regression?

- A single neuron is a logistic regressor!



# What do neural nets have to do with logistic regression?


- A single neuron is a logistic regressor!






# What do neural nets have to do with logistic regression?

- A single neuron is a logistic regressor!



# What do neural nets have to do with logistic regression?


- A single neuron is a logistic regressor!



Multiple inputs: summation of those inputs' weights and the neuron's bias

# What do neural nets have to do with logistic regression?

- A single neuron is a logistic regressor!




$$b + \sum_{i=1}^m w_i \cdot x_i$$



# What do neural nets have to do with logistic regression?

- A single neuron is a logistic regressor!




Or, in linear algebra notation:

$$\begin{bmatrix} 1 & x_1 & x_2 & x_m \end{bmatrix} \begin{bmatrix} b \\ w_1 \\ w_2 \\ w_m \end{bmatrix}$$

# What do neural nets have to do with logistic regression?

- A single neuron is a logistic regressor!

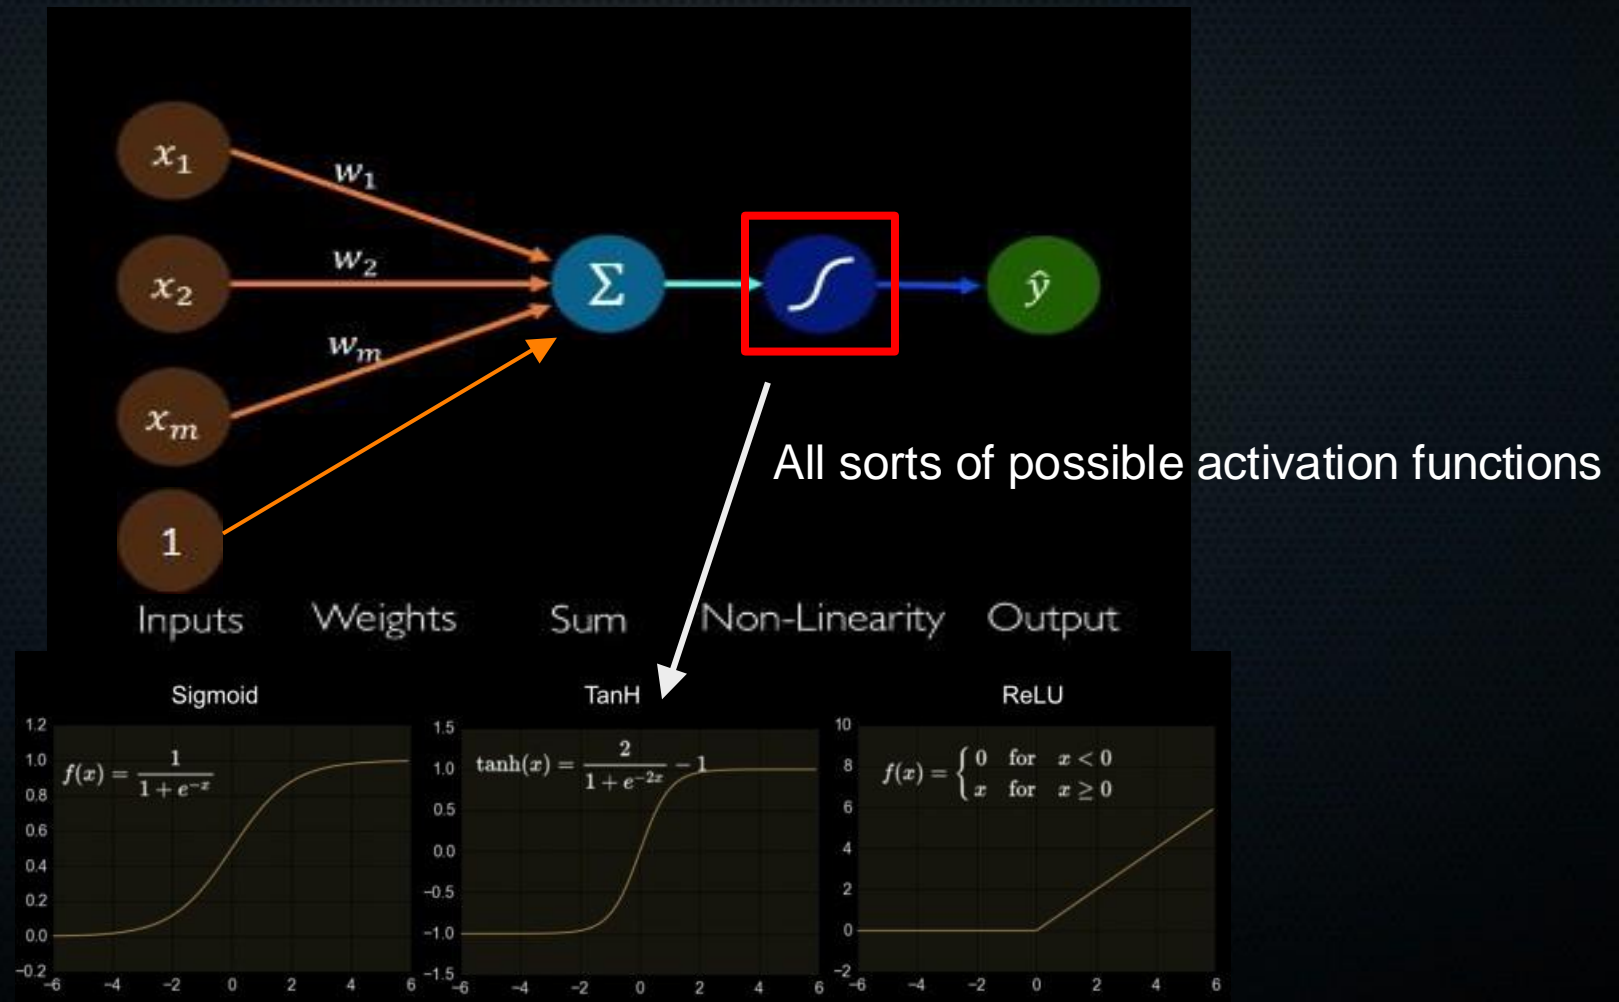


Bias unit (always 1, same for every neuron, most often not shown)

$$\begin{bmatrix} 1 & x_1 & x_2 & x_m \end{bmatrix} \begin{bmatrix} b \\ w_1 \\ w_2 \\ w_m \end{bmatrix}$$

# What do neural nets have to do with logistic regression?

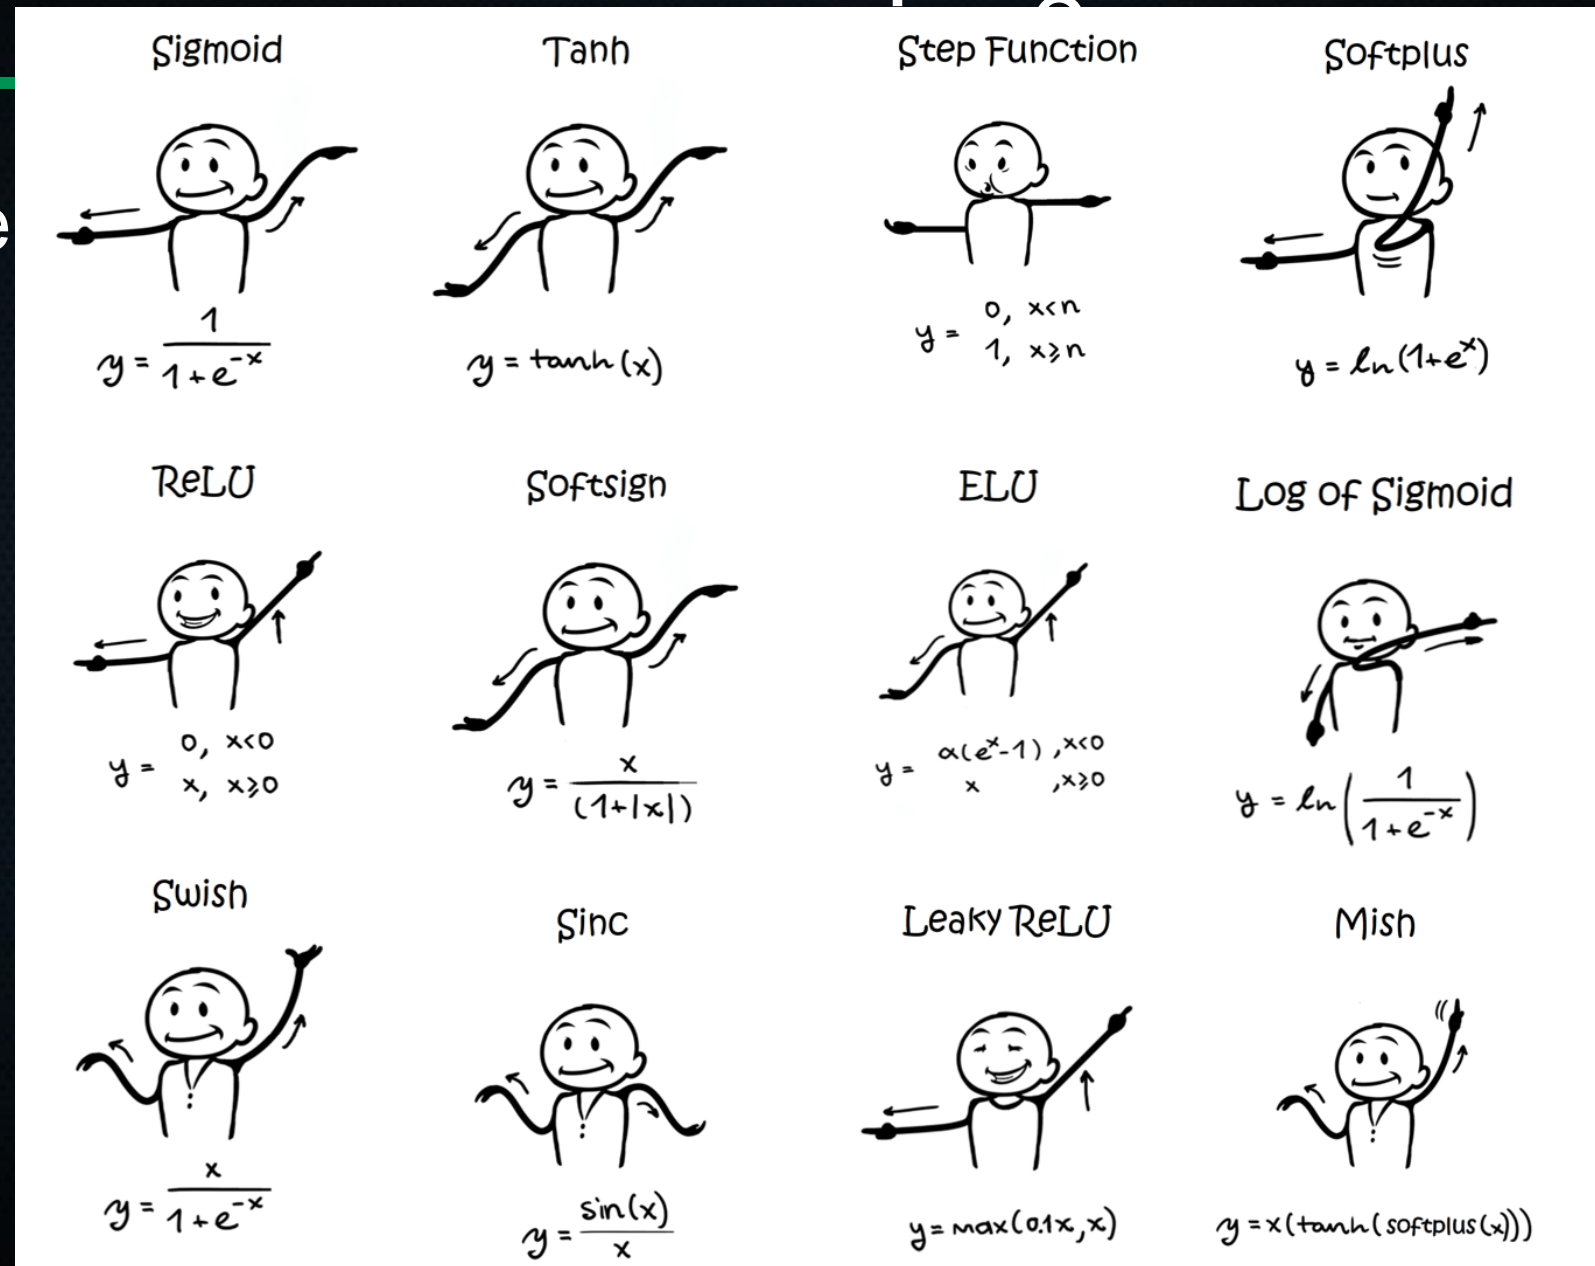
- A single neuron is a logistic regressor!





# What do neural nets have to do with logistic


- A single



n functions

# What do neural nets have to do with logistic regression?

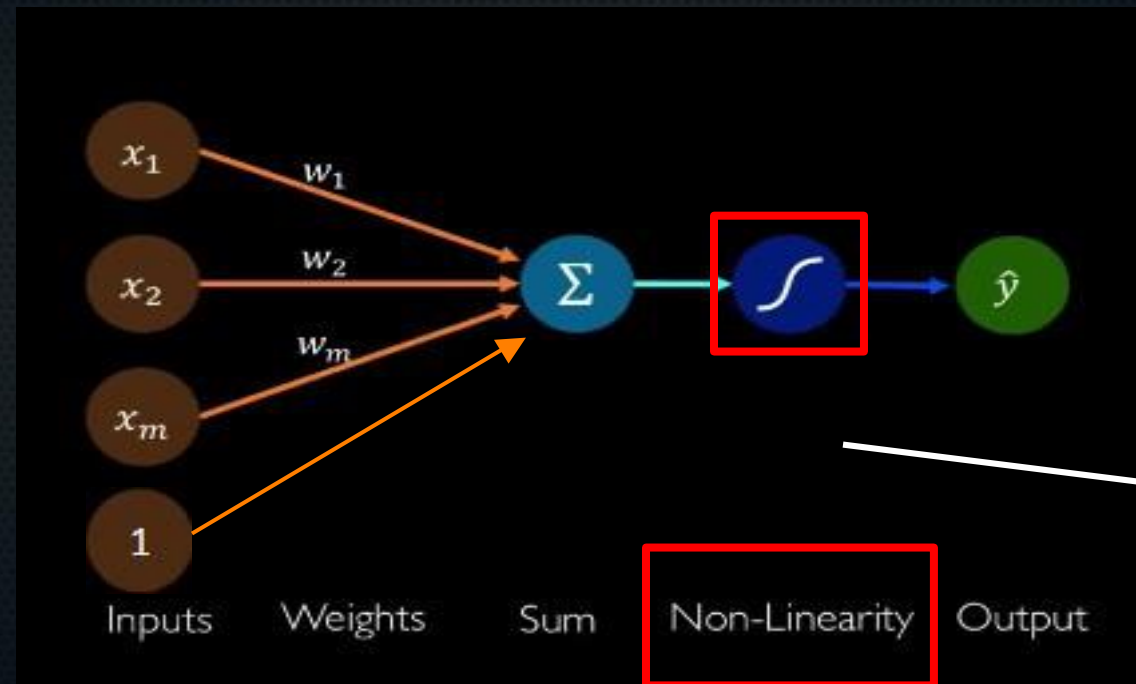
- A single neuron is a logistic regressor!




- Why non-linearity?

# What do neural nets have to do with logistic regression?

- A single neuron is a logistic regressor!

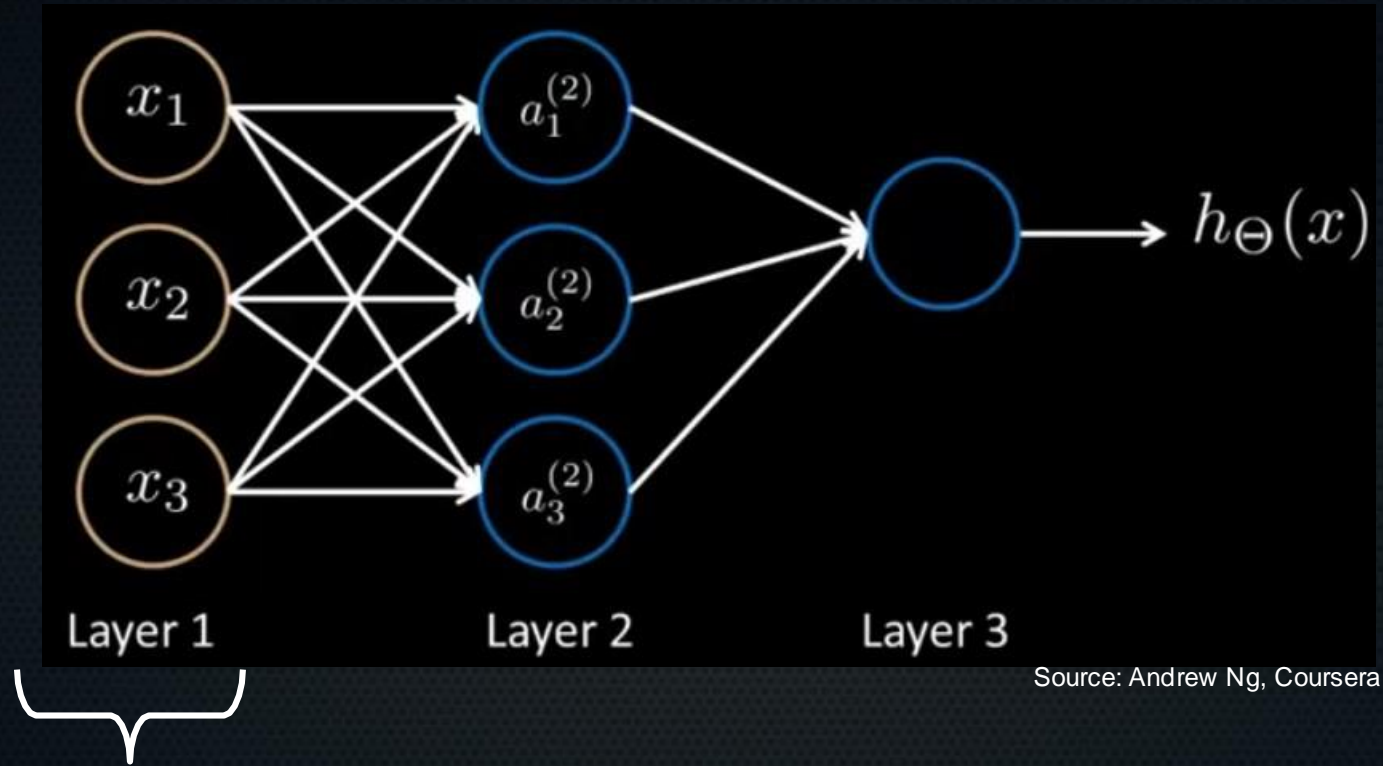


- Why non-linearity? → without them, a NN (no matter how deep) could only approximate linear functions






# Workings of a simple neural network



Input layer: shown as  
neurons, are just features

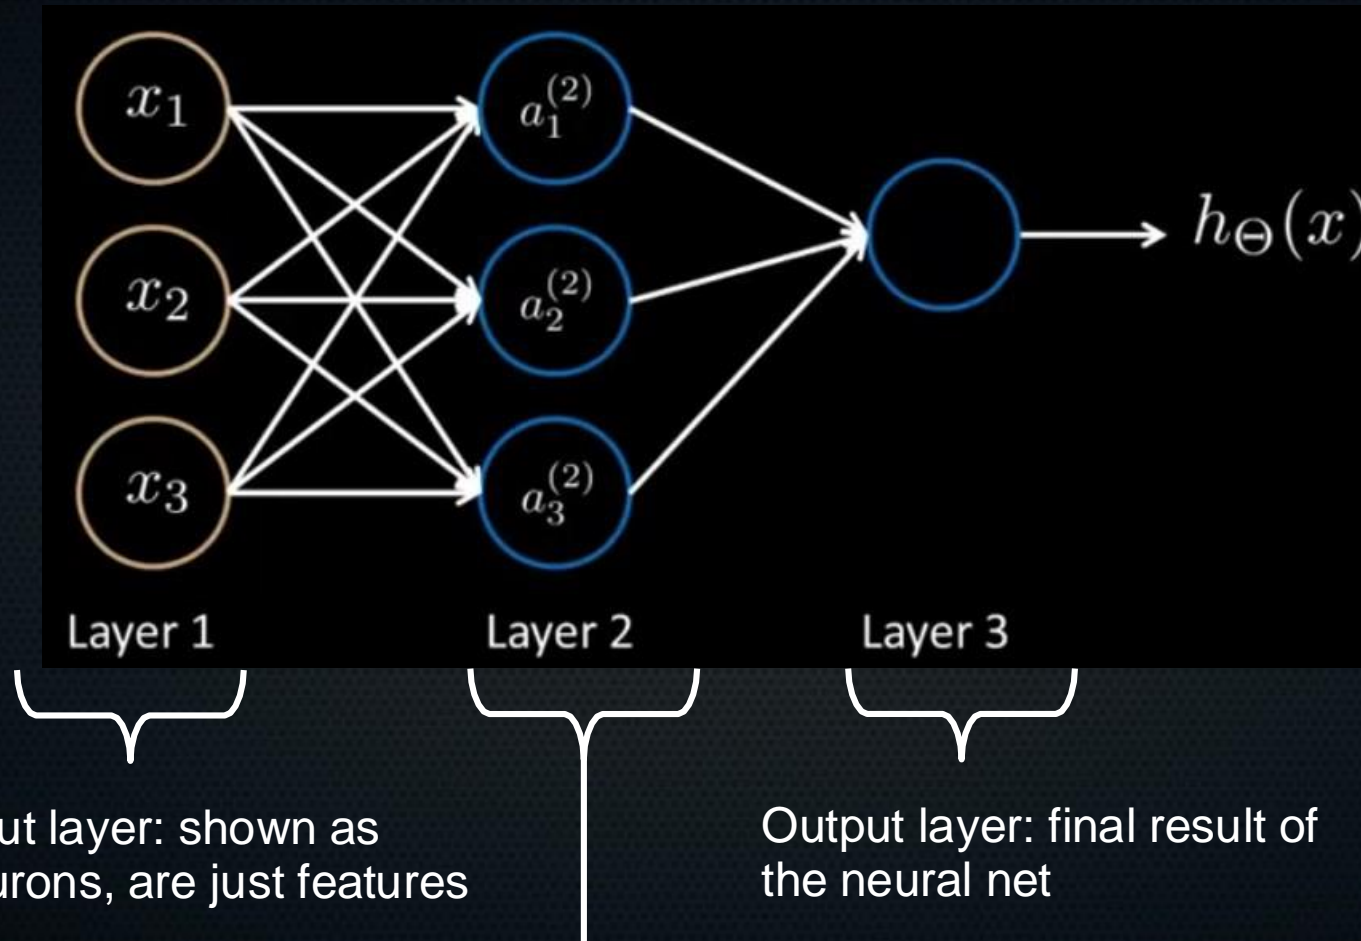
# Workings of a simple neural network



Input layer: shown as neurons, are just features

Output layer: final result of the neural net


# Workings of a simple neural network



Hidden layer(s): intermediate layers whose outputs are not directly observed (hence hidden). Here: 1 HL. Facebook's DenseNet family of NNs had 121-264 HLs in 2016 (0.8-15.3 million parameters).




# Workings of a simple neural network




Activation of neuron 1 in the 2<sup>nd</sup> layer of the network.

# Workings of a simple neural network



Activation of neuron 2 in the 2<sup>nd</sup> layer of the network.


# Workings of a simple neural network



Activation of neuron 1 in the 3<sup>rd</sup> layer of the network.




# What is this network calculating?




$$\sigma \left( \begin{bmatrix} 1 & x_1 & x_2 & x_3 \end{bmatrix} \begin{bmatrix} b \\ w_1 \\ w_2 \\ w_3 \end{bmatrix} \right)$$

# What is this network calculating?



$$\sigma \left( \begin{bmatrix} 1 & x_1 & x_2 & x_3 \end{bmatrix} \begin{bmatrix} b \\ w_1 \\ w_2 \\ w_3 \end{bmatrix} \right) \rightarrow \sigma \left( \begin{bmatrix} 1 & 66 & 39 & 12 \end{bmatrix} \begin{bmatrix} -0.5 \\ -2 \\ 4 \\ 1.5 \end{bmatrix} \right)$$


# What is this network calculating?



$$\sigma \left( \begin{bmatrix} 1 & 66 & 39 & 12 \end{bmatrix} \begin{bmatrix} -0.5 \\ -2 \\ 4 \\ 1.5 \end{bmatrix} \right) \rightarrow \sigma(41.5) \rightarrow 0.999 \dots$$




# What is this network calculating?



$$\sigma \left( \begin{bmatrix} 1 & x_1 & x_2 & x_3 \end{bmatrix} \begin{bmatrix} b \\ w_1 \\ w_2 \\ w_3 \end{bmatrix} \right)$$


# What is this network calculating?



$$\sigma \left( \begin{bmatrix} 1 & x_1 & x_2 & x_3 \end{bmatrix} \begin{bmatrix} b \\ w_1 \\ w_2 \\ w_3 \end{bmatrix} \right)$$

Calculate for all units  
at the same time with  
a theta matrix  $\Theta^{(i)}$

# What is this network calculating?




$$\sigma \left( [1 \quad x_1 \quad x_2 \quad x_3] \begin{bmatrix} b_1 & b_2 & b_3 \\ w_{11} & w_{21} & w_{31} \\ w_{12} & w_{22} & w_{32} \\ w_{13} & w_{23} & w_{33} \end{bmatrix} \right) \rightarrow \sigma ([1 \quad 2 \quad 3])$$

$\downarrow$

$[\sigma(1) \quad \sigma(2) \quad \sigma(3)]$

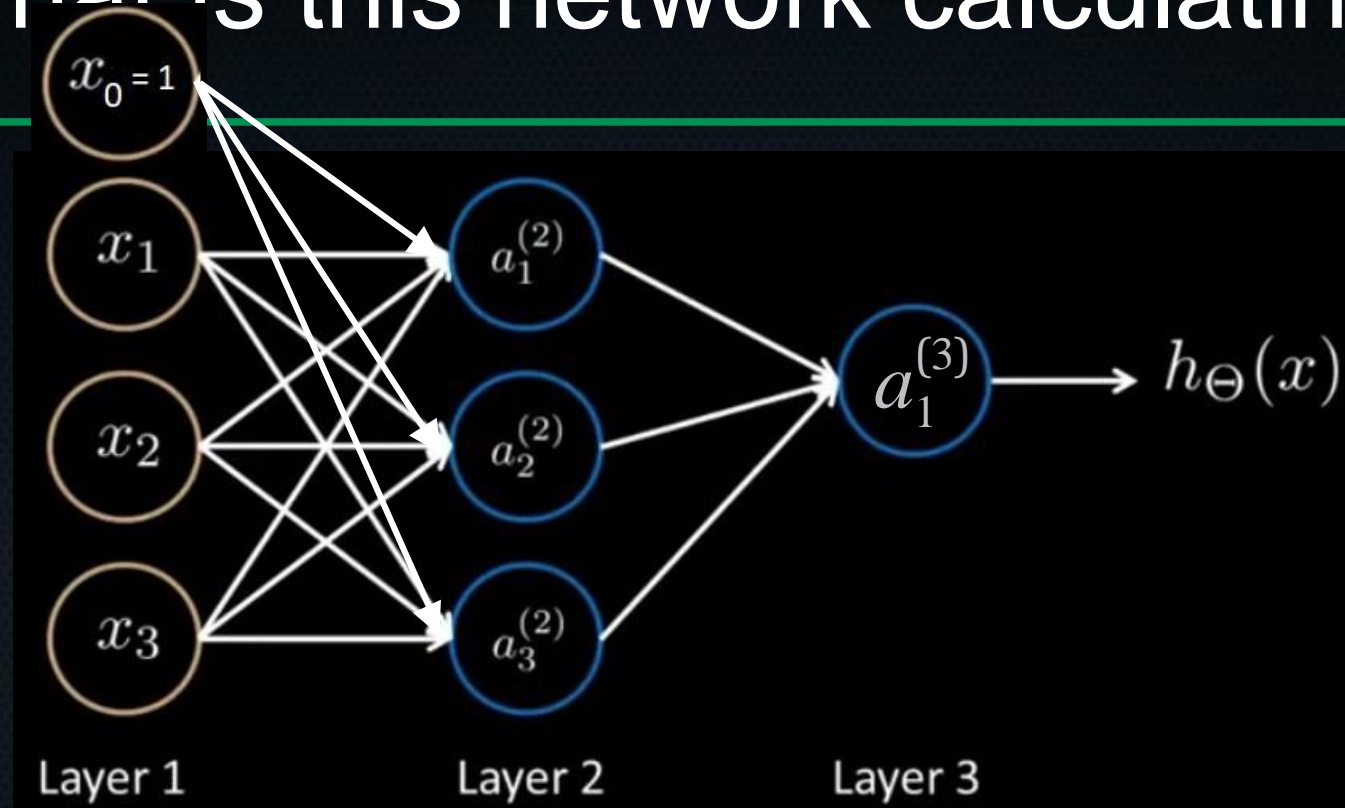


# What is this network calculating?



$$\sigma \left( \begin{bmatrix} b_1 & w_{11} & w_{12} & w_{31} \\ b_2 & w_{21} & w_{22} & w_{23} \\ b_3 & w_{31} & w_{32} & w_{33} \end{bmatrix} \begin{bmatrix} 1 \\ x_1 \\ x_2 \\ x_3 \end{bmatrix} \right) \rightarrow \sigma \left( \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix} \right) \rightarrow \begin{bmatrix} \sigma(1) \\ \sigma(2) \\ \sigma(3) \end{bmatrix}$$

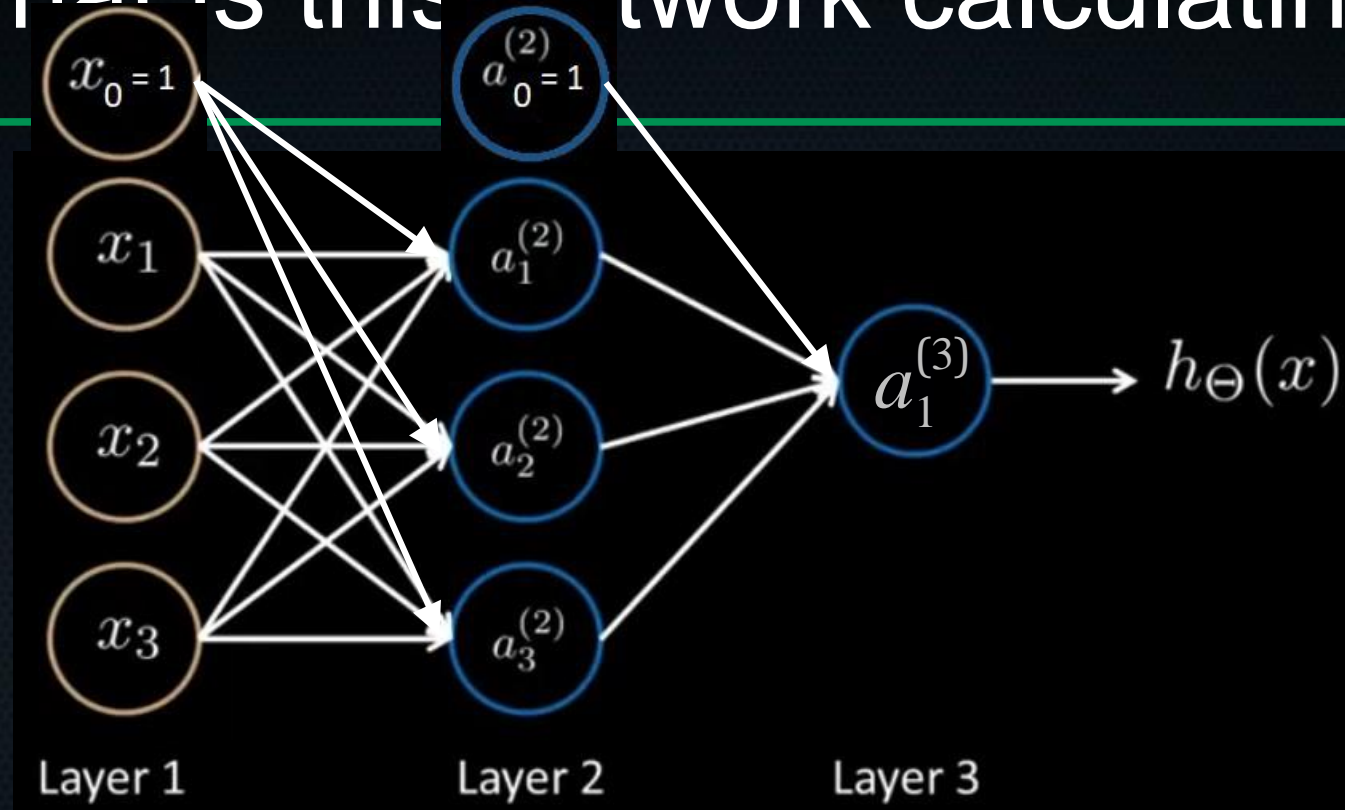
# What is this network calculating?



$$\begin{aligned}a_1^{(2)} &= g(\Theta_{10}^{(1)} x_0 + \Theta_{11}^{(1)} x_1 + \Theta_{12}^{(1)} x_2 + \Theta_{13}^{(1)} x_3) \\a_2^{(2)} &= g(\Theta_{20}^{(1)} x_0 + \Theta_{21}^{(1)} x_1 + \Theta_{22}^{(1)} x_2 + \Theta_{23}^{(1)} x_3) \\a_3^{(2)} &= g(\Theta_{30}^{(1)} x_0 + \Theta_{31}^{(1)} x_1 + \Theta_{32}^{(1)} x_2 + \Theta_{33}^{(1)} x_3)\end{aligned}$$

$$\Theta^{(1)} \text{ (layer 1 to layer 2)}$$
$$\begin{bmatrix} b_1 & w_{11} & w_{12} & w_{13} \\ b_2 & w_{21} & w_{22} & w_{23} \\ b_3 & w_{31} & w_{32} & w_{33} \end{bmatrix}$$

# What is this network calculating?



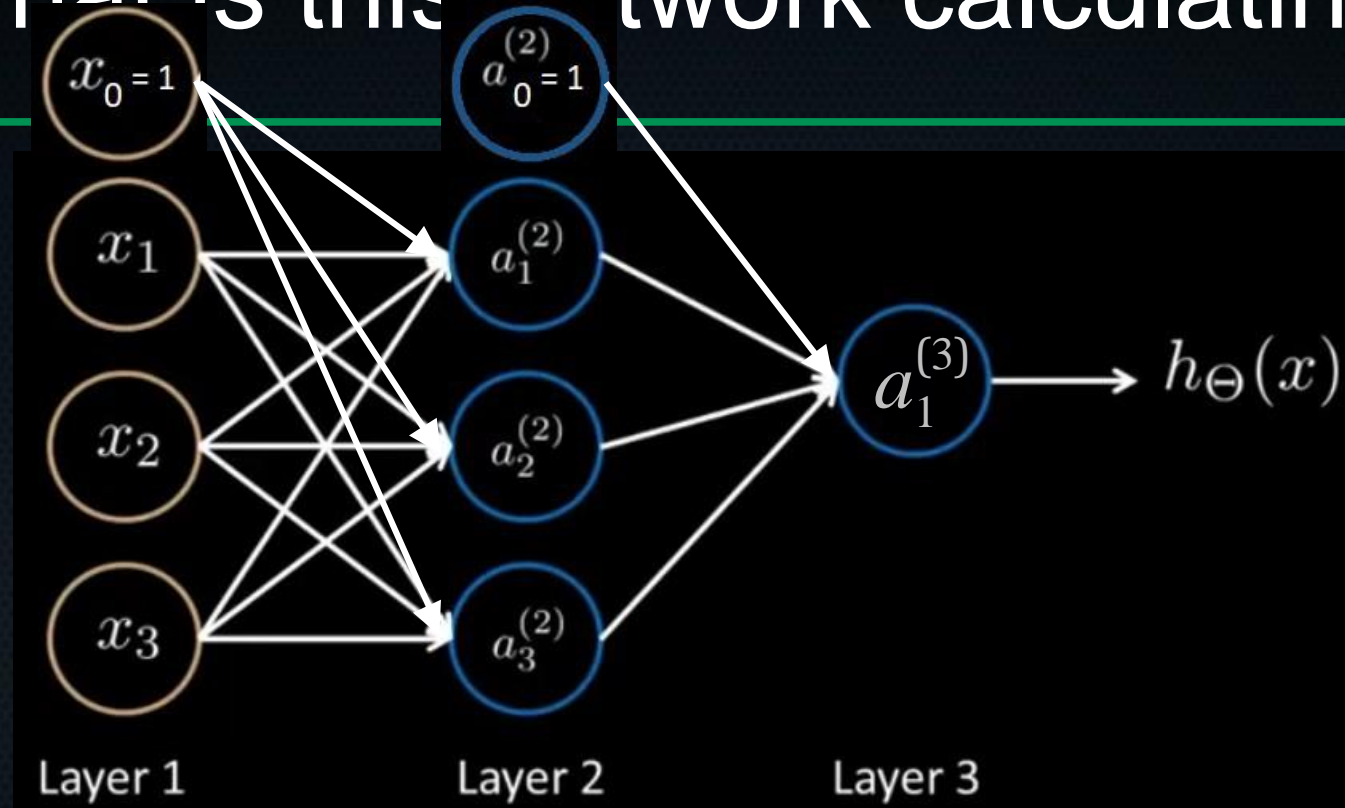
$$h_{\Theta}(x) = a_1^{(3)} = g(\Theta_{10}^{(2)} a_0^{(2)} + \Theta_{11}^{(2)} a_1^{(2)} + \Theta_{12}^{(2)} a_2^{(2)} + \Theta_{13}^{(2)} a_3^{(2)})$$

$\Theta^{(2)}$  (layer 2 to layer 3)

$$\begin{bmatrix} b_1 & w_{11} & w_{12} & w_{13} \end{bmatrix}$$



# What is this network calculating?



- This calculation of the output of the network is called forward propagation

# How do we perform?

---

- Just like before, there is a cost function.
- But we will talk about that and its implementation tomorrow!

# How do we get parameters?

---

- Just like before, there is a cost function and a way to minimise this. But it's a bit more involved.
- To get parameters, we will use the principle of backpropagation. We'll get to that tomorrow.




# Recap so far

---

- Neurons in neural networks are not really like biological neurons, except superficially
- Neural networks can be thought of as hierarchical sets of logistic regressors
- We essentially make earlier layers learn useful features for distinction on their own, and can use these best possible learned features for the classification by the final unit(s)
- Parsing an example through the network and getting the output is called forward propagation
- *Universal approximation* holds that, in principle, neural networks can learn any continuous function arbitrarily well

# Multiclass classification in neural nets




$$h_{\Theta}(x) \in \mathbb{R}^4$$

Want  $h_{\Theta}(x) \approx \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}$ ,  $h_{\Theta}(x) \approx \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}$ ,  $h_{\Theta}(x) \approx \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}$ , etc.  
when pedestrian      when car      when motorcycle



# Multiclass classification in neural nets



$$h_{\Theta}(x) \in \mathbb{R}^4$$

Want  $h_{\Theta}(x) \approx \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}$ ,  $h_{\Theta}(x) \approx \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}$ ,  $h_{\Theta}(x) \approx \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}$ , etc.  
when pedestrian      when car      when motorcycle

Training set:  $(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})$


$y^{(i)}$  one of  $\begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}$ ,  $\begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}$ ,  $\begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}$ ,  $\begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \end{bmatrix}$   
pedestrian   car   motorcycle   truck

~~Previously  
 $y \in \{1, 2, 3, 4\}$~~



# Question to you

- Say we have 10 classes and the following network, how many parameters in  $\Theta^{(2)}$ ?




$\Theta^{(j)}$  = matrix of weights controlling function mapping from layer  $j$  to layer  $j + 1$

$$y^{(10)} = \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

Example:  
10th training  
sample is class 3

# Question to you

- Say we have 10 classes and the following network, how many parameters in  $\Theta^{(2)}$ ?



$\Theta^{(j)}$  = matrix of weights controlling function mapping from layer  $j$  to layer  $j + 1$

$$y^{(10)} = \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$


Example:  
10th training  
sample is class 3

- Answer: 60.  $5 \times 10 \rightarrow$  weights between units.  
+ 10  $\rightarrow$  bias of each unit in output layer



# Question to you

- Say we have 10 classes and the following network, how many parameters in  $\Theta^{(2)}$ ?



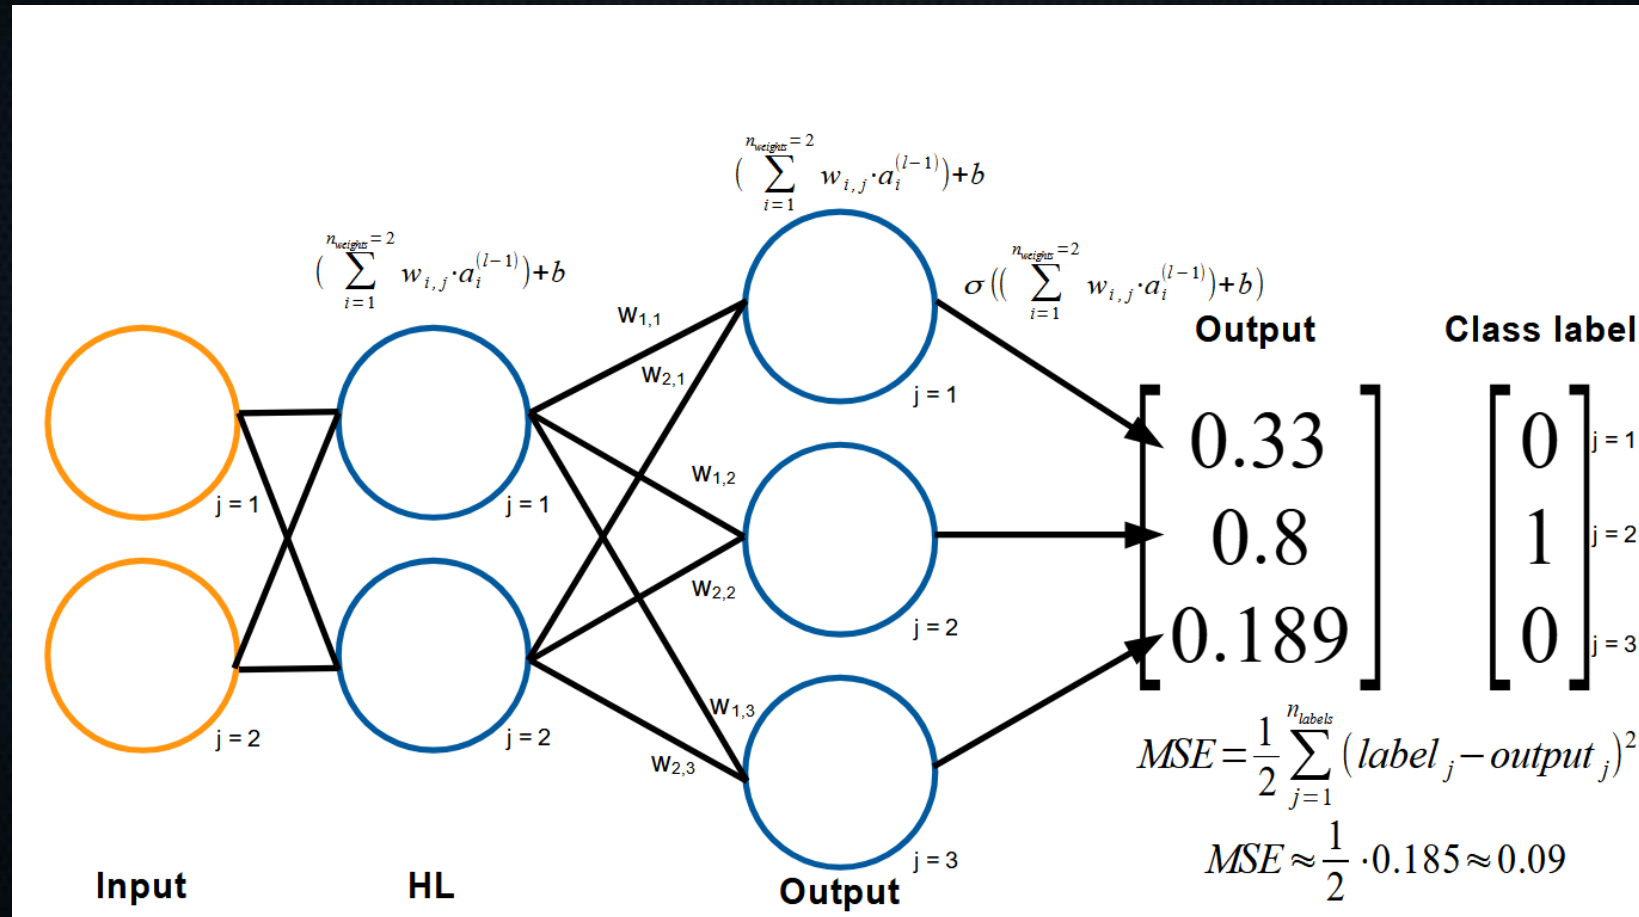
$\Theta^{(2)}$  = matrix of weights controlling function mapping from layer  $j$  to layer  $j + 1$

$$\sigma \left( \begin{bmatrix} b_1 & w_{11} & w_{12} & w_{13} & w_{14} & w_{15} \\ b_2 & w_{21} & w_{22} & w_{23} & w_{24} & w_{25} \\ b_3 & w_{31} & w_{32} & w_{33} & w_{34} & w_{35} \\ b_4 & w_{41} & w_{42} & w_{43} & w_{44} & w_{45} \\ b_5 & w_{51} & w_{52} & w_{53} & w_{54} & w_{55} \end{bmatrix} \begin{bmatrix} 1 \\ a_1^{(2)} \\ a_2^{(2)} \\ a_3^{(2)} \\ a_4^{(2)} \\ a_5^{(2)} \end{bmatrix} \right)$$

- Answer: 60.  $5 \times 10 \rightarrow$  weights between units.  
+ 10  $\rightarrow$  bias of each unit in output layer



# Walk through of linear algebra



# Please git pull

---

To pull changes from git while keeping your own modifications:

```
git stash
```

remove any .ipynb\_checkpoints folder in your practical material if git is mad about them (`rm -rf .ipynb_checkpoints`). These files are saved states of your notebook, but the .ipynb also has a saved state. It's basically a very simple form of backup, which you don't need.

```
git pull
```

```
git pop stash
```

Or use github desktop: <https://desktop.github.com/download/>

# Summary

---

- We can use individual neurons to calculate simple logic functions
- We can combine the outputs of single neurons to calculate more complex (logic) functions
- For multiclass classification, we simply make class a vector, where we strive for the real class to be 1, and all other classes 0.



# Time for the afternoon practical!

\*and continuing where you are

